

pmdarima.arima.auto_arima 🔗

```
pmdarima.arima.auto_arima(y, X=None, start_p=2, d=None, start_q=2, max_p=5, max_d=2, max_q=5,
start_P=1, D=None, start_Q=1, max_P=2, max_D=1, max_Q=2, max_order=5, m=1, seasonal=True,
stationary=False, information_criterion='aic', alpha=0.05, test='kpss', seasonal_test='ocsb', stepwise=True,
n_jobs=1, start_params=None, trend=None, method='lbfgs', maxiter=50, offset_test_args=None,
seasonal_test_args=None, suppress_warnings=True, error_action='trace', trace=False, random=False,
random_state=None, n_fits=10, return_valid_fits=False, out_of_sample_size=0, scoring='mse',
scoring_args=None, with_intercept='auto', sarimax_kwargs=None, **fit_args) \[source\] \[source\]
```

Automatically discover the optimal order for an ARIMA model.

The auto-ARIMA process seeks to identify the most optimal parameters for an `ARIMA` model, settling on a single fitted ARIMA model. This process is based on the commonly-used R function, `forecast::auto.arima` [3].

Auto-ARIMA works by conducting differencing tests (i.e., Kwiatkowski–Phillips–Schmidt–Shin, Augmented Dickey–Fuller or Phillips–Perron) to determine the order of differencing, `d`, and then fitting models within ranges of defined `start_p`, `max_p`, `start_q`, `max_q` ranges. If the `seasonal` optional is enabled, auto-ARIMA also seeks to identify the optimal `P` and `Q` hyper- parameters after conducting the Canova-Hansen to determine the optimal order of seasonal differencing, `D`.

In order to find the best model, auto-ARIMA optimizes for a given `information_criterion`, one of ('aic', 'aicc', 'bic', 'hqic', 'oob') (Akaike Information Criterion, Corrected Akaike Information Criterion, Bayesian Information Criterion, Hannan-Quinn Information Criterion, or “out of bag” –for validation scoring–respectively) and returns the ARIMA which minimizes the value.

Note that due to stationarity issues, auto-ARIMA might not find a suitable model that will converge. If this is the case, a `ValueError` will be thrown suggesting stationarity-inducing measures be taken prior to re-fitting or that a new range of `order` values be selected. Non-stepwise (i.e., essentially a grid search) selection can be slow, especially for seasonal data. Stepwise algorithm is outlined in Hyndman and Khandakar (2008).

Parameters: `y` : array-like or iterable, shape=(`n_samples`),

The time-series to which to fit the `ARIMA` estimator. This may either be a Pandas `Series` object (statsmodels can internally use the dates in the index), or a numpy array. This should be a one-dimensional array of floats, and should not contain any `np.nan` or `np.inf` values.

`X` : array-like, shape=[`n_obs`, `n_vars`], optional (default=None)

An optional 2-d array of exogenous variables. If provided, these variables are used as additional features in the regression operation. This should not include a constant or trend. Note that if an `ARIMA` is fit on exogenous features, it must be provided exogenous features for making predictions.

`start_p` : int, optional (default=2)

The starting value of `p`, the order (or number of time lags) of the auto-regressive (“AR”) model. Must be a positive integer.

`d` : int, optional (default=None)

The order of first-differencing. If None (by default), the value will automatically be selected based on the results of the `test` (i.e., either the Kwiatkowski–Phillips–Schmidt–Shin, Augmented Dickey–Fuller or the Phillips–Perron test will be conducted to find the most probable value). Must be a positive integer or None. Note that if `d` is None, the runtime could be significantly longer.

`start_q` : int, optional (default=2)

The starting value of `q`, the order of the moving-average (“MA”) model. Must be a positive integer.

`max_p` : int, optional (default=5)

The maximum value of `p`, inclusive. Must be a positive integer greater than or equal to `start_p`.

`max_d` : int, optional (default=2)

The maximum value of `d`, or the maximum number of non-seasonal differences. Must be a positive integer greater than or equal to `d`.

`max_q` : int, optional (default=5)

The maximum value of `q`, inclusive. Must be a positive integer greater than `start_q`.

start_P : int, optional (default=1)

The starting value of `p`, the order of the auto-regressive portion of the seasonal model.

D : int, optional (default=None)

The order of the seasonal differencing. If None (by default, the value will automatically be selected based on the results of the `seasonal_test`). Must be a positive integer or None.

start_Q : int, optional (default=1)

The starting value of `q`, the order of the moving-average portion of the seasonal model.

max_P : int, optional (default=2)

The maximum value of `p`, inclusive. Must be a positive integer greater than `start_P`.

max_D : int, optional (default=1)

The maximum value of `D`. Must be a positive integer greater than `D`.

max_Q : int, optional (default=2)

The maximum value of `q`, inclusive. Must be a positive integer greater than `start_Q`.

max_order : int, optional (default=5)

Maximum value of $p+q+P+Q$ if model selection is not stepwise. If the sum of `p` and `q` is \geq `max_order`, a model will *not* be fit with those parameters, but will progress to the next combination. Default is 5. If `max_order` is None, it means there are no constraints on maximum order.

m : int, optional (default=1)

The period for seasonal differencing, `m` refers to the number of periods in each season. For example, `m` is 4 for quarterly data, 12 for monthly data, or 1 for annual (non-seasonal) data. Default is 1. Note that if `m` == 1 (i.e., is non-seasonal), `seasonal` will be set to False. For more information on setting this parameter, see [Setting m](#).

seasonal : bool, optional (default=True)

Whether to fit a seasonal ARIMA. Default is True. Note that if `seasonal` is True and `m` == 1, `seasonal` will be set to False.

stationary : bool, optional (default=False)

Whether the time-series is stationary and `d` should be set to zero.

information_criterion : str, optional (default='aic')

The information criterion used to select the best ARIMA model. One of `pmdarima.arima.auto_arima.VALID_CRITERIA`, ('aic', 'bic', 'hqic', 'oob').

alpha : float, optional (default=0.05)

Level of the test for testing significance.

test : str, optional (default='kpss')

Type of unit root test to use in order to detect stationarity if `stationary` is False and `d` is None. Default is 'kpss' (Kwiatkowski–Phillips–Schmidt–Shin).

seasonal_test : str, optional (default='ocsb')

This determines which seasonal unit root test is used if `seasonal` is True and `d` is None. Default is 'OCSB'.

stepwise : bool, optional (default=True)

Whether to use the stepwise algorithm outlined in Hyndman and Khandakar (2008) to identify the optimal model parameters. The stepwise algorithm can be significantly faster than fitting all (or a `random` subset of) hyper-parameter combinations and is less likely to over-fit the model.

n_jobs : int, optional (default=1)

The number of models to fit in parallel in the case of a grid search (`stepwise=False`). Default is 1, but -1 can be used to designate “as many as possible”.

start_params : array-like, optional (default=None)

Starting parameters for `ARMA(p,q)`. If None, the default is given by `ARMA._fit_start_params`.

method : str, optional (default='lbfgs')

The `method` determines which solver from `scipy.optimize` is used, and it can be chosen from among the following strings:

- 'newton' for Newton-Raphson
- 'nm' for Nelder-Mead
- 'bfgs' for Broyden-Fletcher-Goldfarb-Shanno (BFGS)
- 'lbfgs' for limited-memory BFGS with optional box constraints
- 'powell' for modified Powell's method
- 'cg' for conjugate gradient
- 'nCG' for Newton-conjugate gradient
- 'basinhopping' for global basin-hopping solver

The explicit arguments in `fit` are passed to the solver, with the exception of the basin-hopping solver. Each solver has several optional arguments that are not the same across solvers. These can be passed as `**fit_kwargs`

trend : str or None, optional (default=None)

The trend parameter. If `with_intercept` is True, `trend` will be used. If `with_intercept` is False, the trend will be set to a no- intercept value.

maxiter : int, optional (default=50)

The maximum number of function evaluations. Default is 50.

offset_test_args : dict, optional (default=None)

The args to pass to the constructor of the offset (`d`) test. See `pmdarima.arima.stationarity` for more details.

seasonal_test_args : dict, optional (default=None)

The args to pass to the constructor of the seasonal offset (`D`) test. See `pmdarima.arima.seasonality` for more details. Examples of valid kwargs will vary based on the test. For the `OCSBTest` (default) they include:

- 'lag_method'
- 'max_lag'

suppress_warnings : bool, optional (default=True)

Many warnings might be thrown inside of statsmodels. If `suppress_warnings` is True, all of the warnings coming from `ARIMA` will be squelched. Note that this will not suppress UserWarnings created by bad argument combinations.

error_action : str, optional (default='warn')

If unable to fit an `ARIMA` for whatever reason, this controls the error-handling behavior. Model fits can fail for linear algebra errors, convergence errors, or any number of problems related to stationarity or input data.

- 'warn': Warns when an error is encountered (default)
- 'raise': Raises when an error is encountered
- 'ignore': Ignores errors (not recommended)
- **'trace': Logs the entire error stacktrace and continues the**
search. This is the best option when trying to determine why a model is failing.

trace : bool or int, optional (default=False)

Whether to print status on the fits. A value of False will print no debugging information. A value of True will print some. Integer values exceeding 1 will print increasing amounts of debug information at each fit.

random : bool, optional (default=False)

Similar to grid searches, `auto_arima` provides the capability to perform a “random search” over a hyper-parameter space. If `random` is True, rather than perform an exhaustive search or `stepwise` search, only `n_fits` ARIMA models will be fit (`stepwise` must be False for this option to do anything).

random_state : int, long or numpy `RandomState`, optional (default=None)

The PRNG for when `random=True`. Ensures replicable testing and results.

n_fits : int, optional (default=10)

If `random` is True and a “random search” is going to be performed, `n_fits` is the number of ARIMA models to be fit.

return_valid_fits : bool, optional (default=False)

If True, will return all valid ARIMA fits in a list. If False (by default), will only return the best fit.

out_of_sample_size : int, optional (default=0)

The `ARIMA` class can fit only a portion of the data if specified, in order to retain an “out of bag” sample score. This is the number of examples from the tail of the time series to hold out and use as validation examples. The model will not be fit on these samples, but the observations will be added into the model’s `endog` and `exog` arrays so that future forecast values originate from the end of the endogenous vector.

For instance:

```
y = [0, 1, 2, 3, 4, 5, 6]
out_of_sample_size = 2

> Fit on: [0, 1, 2, 3, 4]
> Score on: [5, 6]
> Append [5, 6] to end of self.arma_res_.data.endog values
```

scoring : str, optional (default='mse')

If performing validation (i.e., if `out_of_sample_size` > 0), the metric to use for scoring the out-of-sample data. One of ('mse', 'mae')

scoring_args : dict, optional (default=None)

A dictionary of key-word arguments to be passed to the `scoring` metric.

with_intercept : bool or str, optional (default="auto")

Whether to include an intercept term. Default is “auto” which behaves like True until a point in the search where the sum of differencing terms will explicitly set it to True or False.

sarimax_kwargs : dict or None, optional (default=None)

Keyword arguments to pass to the ARIMA constructor.

****fit_args** : dict, optional (default=None)

A dictionary of keyword arguments to pass to the `ARIMA.fit()` method.

! See also

`pmdarima.arma.ARIMA()`

Notes

- Fitting with *stepwise=False* can prove slower, especially when *seasonal=True*.

References

[R68] https://wikipedia.org/wiki/Autoregressive_integrated_moving_average

[R69] R's auto-arma source code: <https://github.com/robjhyndman/forecast/blob/master/R/arma.R>

[R70] R's auto-arma documentation: <https://www.rdocumentation.org/packages/forecast>